

Neural Network Documentation

Dataset

The dataset used consists of randomly generated features and labels:

- **Features (X):** Random data with 300 samples and 2 features per sample.
- **Labels (y):** Random integers between 0 and 2, indicating 3 classes.

Model Architecture

The neural network architecture is as follows:

- **Input Layer:** 2 features
- **Hidden Layer:** 25 neurons with ReLU activation
- **Output Layer:** Softmax layer with 3 neurons (corresponding to the 3 classes)

Training Process

The neural network is trained using stochastic gradient descent with the following parameters:

- **Learning Rate:** 0.1
- **Regularization Parameter (λ):** 0.001
- **Number of Iterations:** 10,000

During training:

- Forward propagation computes predictions (probs).
- Cross-entropy loss is used to compute the cost function.
- Backpropagation computes gradients and updates parameters ($W1, b1, W2, b2$).
- Cost is printed every 1000 iterations for monitoring.

Accuracy and Insights

Accuracy is typically evaluated using a separate validation or test set:

- **Validation/Test Set:** Split data into training and validation/test sets.
- **Prediction:** Use trained parameters to predict outputs.
- **Accuracy Calculation:** Compare predicted outputs with actual labels.




Figure 1: Neural Network Architecture

In this image the boundaries are decision boundaries each colour represents different class